

DR. JENNIFER WANG

Machine Learning Engineer

Email: jennifer.wang@email.com | Phone: (650) 555-0177

Location: Mountain View, CA | LinkedIn: [linkedin.com/in/jenniferwang](https://www.linkedin.com/in/jenniferwang)

Google Scholar: scholar.google.com/jennifer-wang

PROFESSIONAL SUMMARY

Ph.D. in Machine Learning with 8 years of experience building and deploying production ML systems. Specialized in NLP, computer vision, and recommendation systems at scale. Published 12 peer-reviewed papers and 3 patents.

TECHNICAL EXPERTISE

Machine Learning: Deep Learning, NLP, Computer Vision, Reinforcement Learning, Transfer Learning

Frameworks: PyTorch, TensorFlow, scikit-learn, Hugging Face Transformers, OpenCV

Languages: Python, C++, SQL, R

MLOps: MLflow, Kubeflow, SageMaker, Docker, Kubernetes

Cloud: AWS (SageMaker, EC2, S3), GCP (Vertex AI, BigQuery)

Data: Pandas, NumPy, Spark, Airflow

PROFESSIONAL EXPERIENCE

Senior Machine Learning Engineer | DataDrive AI | Mountain View, CA

April 2020 - Present

- Lead ML team of 5 engineers building recommendation system serving 10M+ users
- Designed and deployed transformer-based NLP model improving search relevance by 35%
- Built end-to-end ML pipeline with automated retraining reducing model staleness from 2 weeks to 2 days
- Implemented A/B testing framework measuring model impact on user engagement (+18% CTR)
- Optimized model inference reducing latency from 200ms to 45ms using model quantization and TensorRT
- Mentored junior ML engineers and conducted technical interviews

Machine Learning Engineer | AI Research Labs | Palo Alto, CA

June 2017 - March 2020

- Developed computer vision models for object detection achieving 92% mAP on custom dataset
- Built data preprocessing pipelines processing 500GB+ of image data daily using Spark

- Deployed ML models to production using Docker and Kubernetes
- Collaborated with research team publishing 4 papers at top-tier ML conferences (NeurIPS, ICML)
- Implemented MLOps best practices including model versioning, monitoring, and rollback strategies

Research Scientist | University ML Lab | Stanford, CA

September 2015 - May 2017

- Conducted PhD research on deep reinforcement learning for robotics
- Published 6 papers in top conferences and journals
- Taught graduate-level machine learning course (4.8/5.0 rating)
- Supervised 3 master's thesis projects

EDUCATION

Ph.D. in Computer Science (Machine Learning) | Stanford University | 2017

Dissertation: "Deep Reinforcement Learning for Robotic Manipulation"

Advisor: Prof. Andrew Ng

M.S. in Computer Science | Stanford University | 2015

B.S. in Computer Science and Mathematics | MIT | 2013

Magna Cum Laude, GPA: 3.9/4.0

PUBLICATIONS & PATENTS

- Wang, J. et al. "Efficient Transformers for NLP at Scale" - NeurIPS 2023
- Wang, J. et al. "Few-Shot Learning for Recommendation Systems" - ICML 2022
- 3 US Patents on ML model optimization and deployment

AWARDS & HONORS

- "40 Under 40 in AI" - AI Magazine (2023)
- Best Paper Award - ICML 2022
- Google PhD Fellowship (2015-2017)

CERTIFICATIONS

- AWS Certified Machine Learning - Specialty (2023)
- TensorFlow Developer Certificate (2021)